NGUYỄN MINH PHƯƠNG

AI Engineer Intern

phuongnm150505@gmail.com | +84 896 980 529 | Ho Chi Minh City | github.com/minhphuong150505 | linkedin.com/in/phuongnm150505

TECHNICAL SKILLS

Programming: Python, Java, SQL, TypeScript, JavaScript
AI / ML: PyTorch, Hugging Face Transformers, Transformer Architecture, LLM Fine-tuning (QLoRA / PEFT), Vision-Language Models, Multi-label Image Classification, OpenCV, torchvision, Grad-CAM, scikit-learn, NumPy, pandas, Matplotlib
LLM & RAG: RAG, Embeddings, Vector Database (Milvus), LangChain text splitters, Prompt Engineering, MedGemma, Ollama
Evaluation & Benchmarking: AUROC, AUPRC, F1, Precision / Recall, Threshold Calibration, Bootstrap CIs, Ablation Studies, BLEU, ROUGE-L, BERTScore
Backend & Deployment: FastAPI, Spring Boot 3, REST API, PostgreSQL, Docker, Docker Compose, GCP, Linux, CUDA
Tools: Git, GitHub, pytest, W&B, Jira, AI Agents (Claude Code, Codex, GitHub Copilot)

KEY PROJECTS

META-CXR - Multimodal AI for Chest X-ray Report Generation | Graduation thesis • Solo GitHub

- Reproduced and extended a two-stage vision-language model (IEEE Access 2025) in PyTorch on MIMIC-CXR: two medical vision encoders (BioViL-T, PubMedCLIP) + Q-Former classify **14** abnormalities, then MedGemma-4B fine-tuned with QLoRA writes the report.
- Model optimization: partially unfreezing both encoders raised test macro-AUROC **0.744** → **0.764** on **3,269** held-out studies (paired bootstrap CI above 0, gains on **14/14** labels), with thresholds calibrated on validation only.
- Fixed an image-representation bug: restored each encoder's native patch grid and removed a constant component by per-image centering, cutting pairwise patch cosine similarity from **0.674** to **-0.014** at only **+0.6%** step time.
- Reproducible experiments: controlled A/B runs dropped components with no measurable gain (Grad-CAM explanation loss, deeper unfreezing); joining MS-CXR boxes by image ID exposed **166** boxes leaking into the test split; **1,012** pytest tests passed.
- LLM debugging: found the chat stop token was overridden, so **78%** of generated tokens came after the model had finished.
- Infrastructure: full training on one **16 GB** GPU (fixed recurring kernel faults traced to unstable RAM settings); multi-GPU DDP smoke runs on **2× T4** with checkpoints on Google Cloud Storage.

FPT Tri Thuc - Enterprise RAG Assistant | Solo GitHub

- Python ingestion pipeline (PDF / web / text loaders, chunking, embedding, hash dedup) into Milvus 2.4, plus a Spring Boot RAG API for Vietnamese Q&A over FPT corporate documents.
- Pulls **64** candidates from Milvus, re-ranks by cosine plus keyword / bigram coverage, cites sources and abstains when evidence is weak; shows each chunk's cosine score in the UI so retrieval quality can be audited.
- Token streaming over SSE with JSON fallback and per-upstream timeouts; **25** JUnit / Mockito tests and **9** Playwright E2E tests; runs on Docker Compose.

Chinese Learning Platform - Full-Stack AI Web App | Solo GitHub

- Designed, built and deployed the whole product alone in under **2 months**; used by friends, and iterated on their feedback.
- Integrated AI into a real product: a Spring Boot API orchestrates an LLM (chat, translation, writing feedback), Azure Speech and a Python FastAPI service for F0 pitch analysis (per-syllable tone scoring); PostgreSQL + Flyway, Docker, deployed on GCP over HTTPS.

WORK EXPERIENCE

FPT Software | Java Developer Intern Ho Chi Minh City | May 2026 - Oct 2026

Project BUHINH - on-premise to cloud migration of a Japanese automaker's manufacturing BOM (MBOM) system
Tech stack: Java EE, Spring Boot, Spring MVC, JSP, JavaScript, JDBC, SQL, IBM Db2, Maven, GitLab, Jira, GitHub Copilot

- Helped the team finish migration Phase 2 ahead of schedule by moving legacy Java EE sub-systems to Spring Boot / Spring MVC, using GitHub Copilot under the team's prompt rules and hand-reviewing every pattern it flagged.
- Brought the cloud version of the parts-applicability search subsystem back in line with the legacy on-premise system: about 50 UI and 30 logic defects fixed and merged through GitLab merge requests.
- Found about 15 on-premise vs cloud differences, logged them on Jira with reproducible steps and evidence, and resolved about 15 tickets as a full team member, on or ahead of deadline; internship rated 10/10 by the company.

EDUCATION

University of Information Technology (UIT) - VNU-HCM | Information Systems 2023 - 2027
GPA: 8.5/10 | Final-year student, all required coursework completed | Expected graduation: Jun 2027

AWARDS AND CERTIFICATES

- Outstanding Student - UIT - VNU-HCM, Semester 2 of academic year 2024 - 2025 2025
- Academic Scholarship - UIT - VNU-HCM, Top 8% of major, awarded for 2 semesters 2023 - 2025

LANGUAGES

English: B2 (CEFR) equivalent - reads technical documentation and AI/ML research papers with ease